

CREECelerator

Composable Hardware Generators for Disaggregated Storage Systems

Bob Zhou, Kyle Kovacs, Vighnesh Iyer, Tan Nguyen

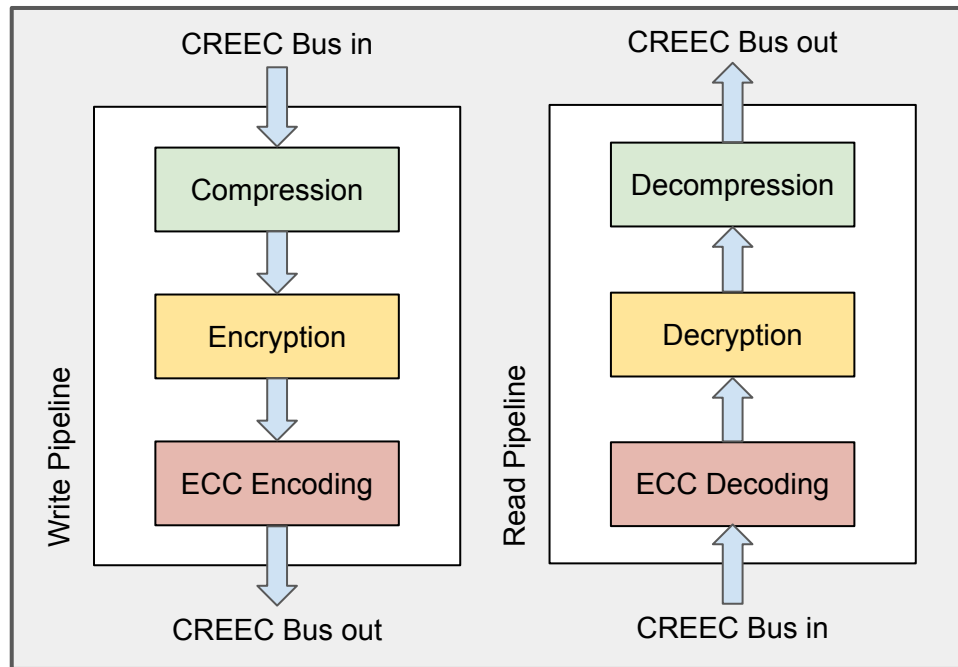
EE290C Fall 2018

Dec 5, 2018

Final Presentation

What We Built

- Composable accelerators for storage in a disaggregated datacenter
- Compression
 - Increases SSD endurance
- Encryption
 - Ensures security of stored data
- ECC
 - Protects data against corruption
- Interconnect
 - Ties everything together with common abstractions and interfaces



Interconnect

- Minimal system bus that can bridge to existing testchipip BlockDevice, or a generic stream bus (AXI4-Stream), or any memory mapped bus (TL)
 - In hindsight, we should have just used AXI4-Stream and a modified architecture

```
case class BusParams(maxBeats: Int, maxInFlight: Int, dataWidth: Int)
```

```
class Header(val p: BusParams) extends Bundle {  
  val len = UInt()  
  val id = UInt()  
}
```

```
class Data(val p: BusParams) extends Bundle {  
  val data = UInt()  
  val id = UInt()  
}
```

```
class CREECBus(val p: BusParams) extends Bundle {  
  val header = Decoupled(Header(p))  
  val data = Decoupled(Data(p))  
}
```

Compression

Problem: SSDs degrade over time

- Endurance is based on number of read/write cycles



Goal: increase SSD lifetime by decreasing total physical I/O cycles

Method: Compress data before it hits the disk

- For the purposes of this project, a simple compression algorithm was used
 - Differential + run-length encoding


Differential Encoding

- Byte sequences are transformed into [initial value, differences]
- Reversible transformation
- Prepares the data for run-length encoding
- Converts any sequence of repeated bytes to a sequence of zeros



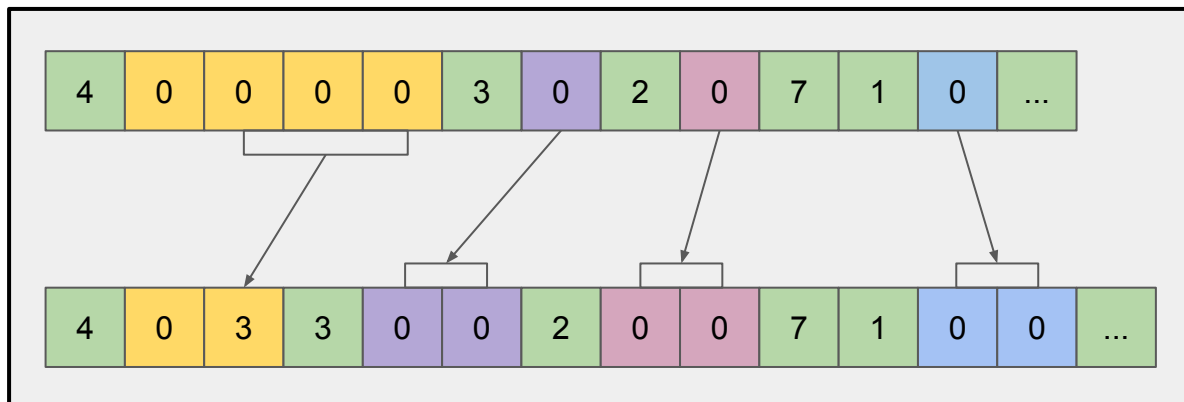
Run-Length Encoding

- “Runs” of zeros are transformed into [0, # additional zeros]
- Reversible transformation
- After differential, compresses runs of up to 256 repeated values into just 3 bytes



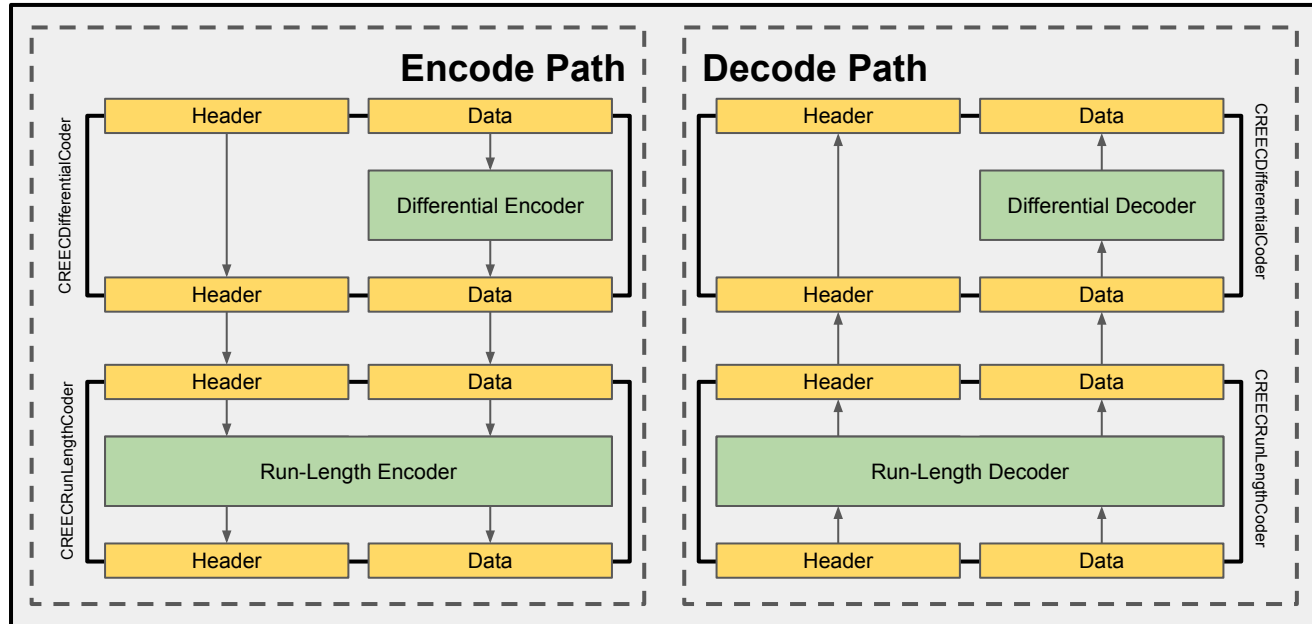
Problem With Run-Length Encoding

- A problem occurs when you have a single zero
 - It expands to 2 bytes
- In practice, this only works well for a specific type of data
 - Data like ASCII text, JPEG images, or other files do not compress well with this method



System Integration

- To interface with the system bus, the block must accept headers and data
- In the case of compression, the output header can only be sent after the data is compressed and the final length is determined



ECC: Reed-Solomon code

RS(n, k, b) **Generator**

- n: number of output symbols
- k: number of input symbols
- b: bitwidth of a symbol

Note: **block-level** parity scheme!

Capacity: correct up to $(n-k)/2$ symbols

Example: RS(16, 8, 8)

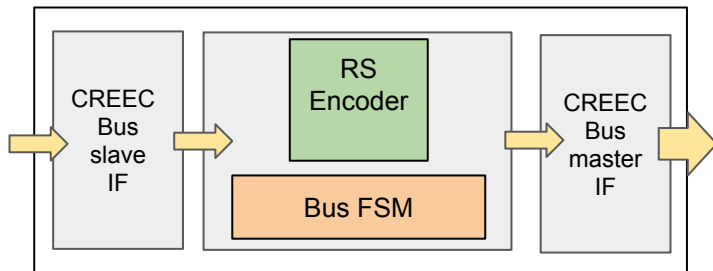
64-bit data

B	B	B	B	B	B	B	B
1	2	3	4	5	6	7	8

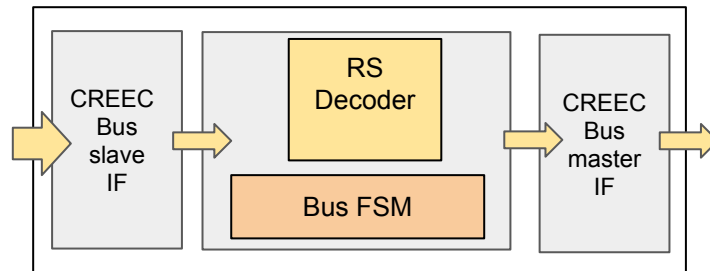


128-bit data

B	B	B	B	B	B	B	B	p	p	p	p	p	p	p	p
1	2	3	4	5	6	7	8	1	2	3	4	5	6	7	8

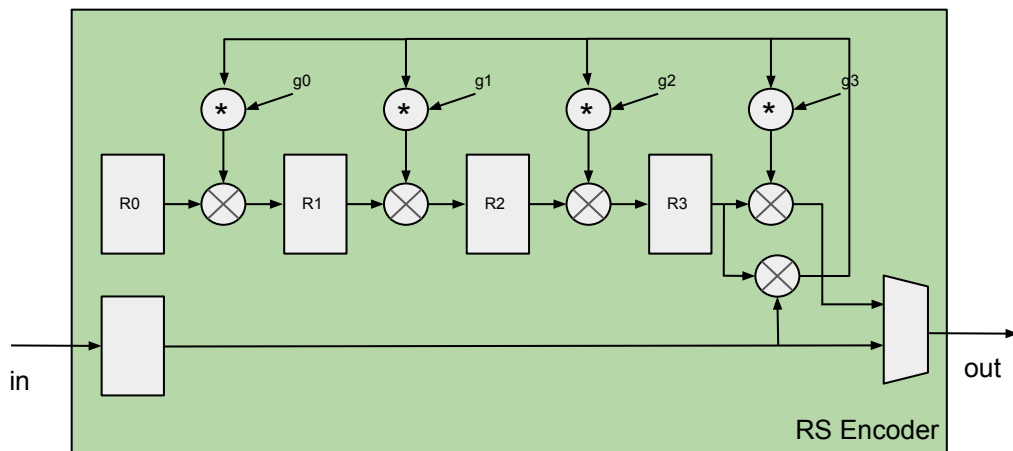


CREEC RS Encoder block



CREEC RS Decoder block

RS Encoder



in₁, in₂, ..., in_k

in₁, in₂, ... in_k,
p₁, p₂, ..., p_(n-k)

Linear-Feedback Shift-Register

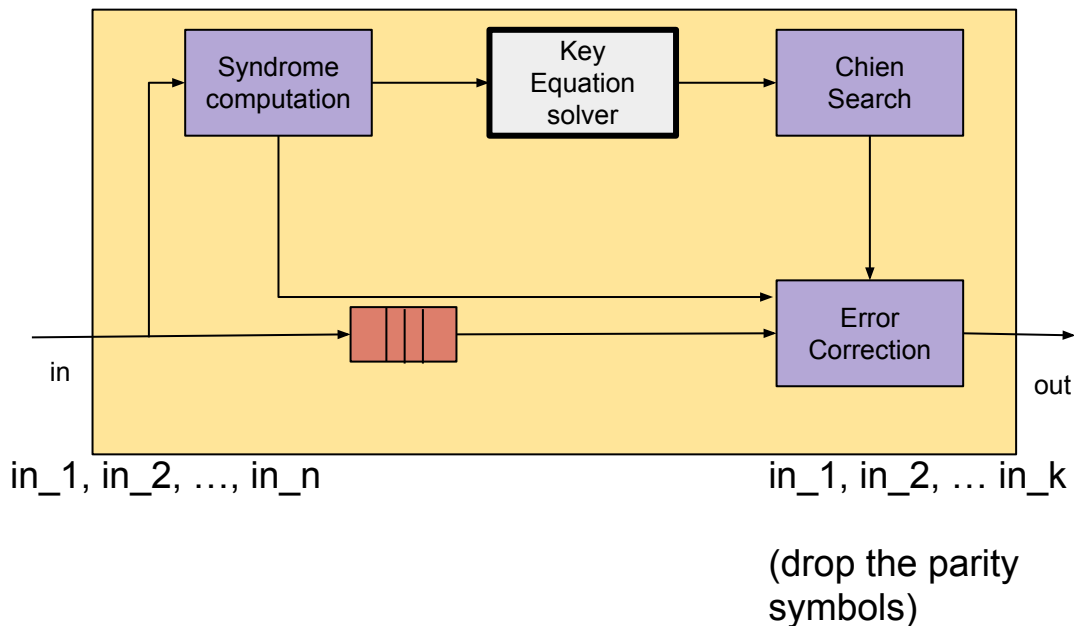
#Registers == #parity symbols

Arithmetic operations are performed in Galois Finite Field

- Addition: XOR operation
- Multiplication: XOR + bitmask + shift operations

No carry generated!

RS Decoder



> 20x expensive than Encoder!

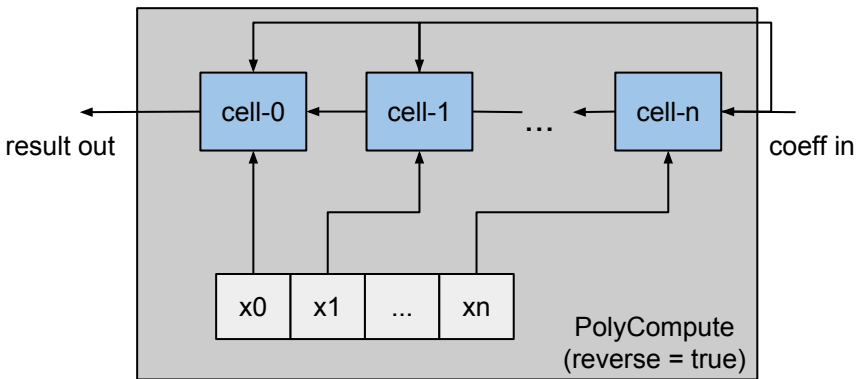
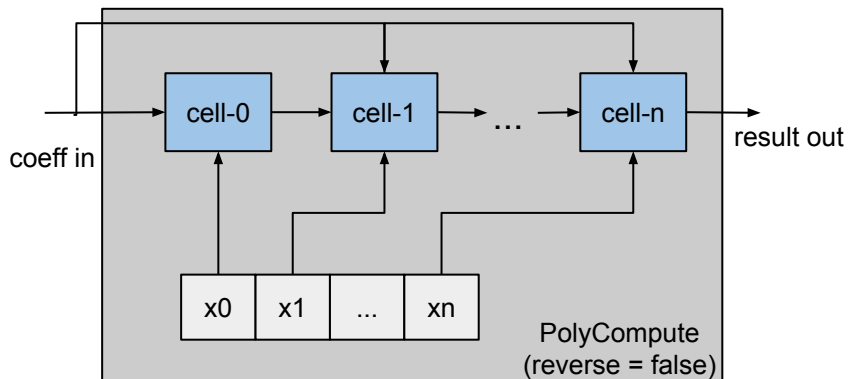
Fixes up to $(n-k)/2$ input symbols

Key Equation Solver: **polynomial division** (Modified Euclidean Algorithm)


Syndrome computation, Chien Search, Error Correction: **polynomial evaluation**

Inversion operation (in Error Correction) is on the critical path

Polynomial evaluation (Horner's method)

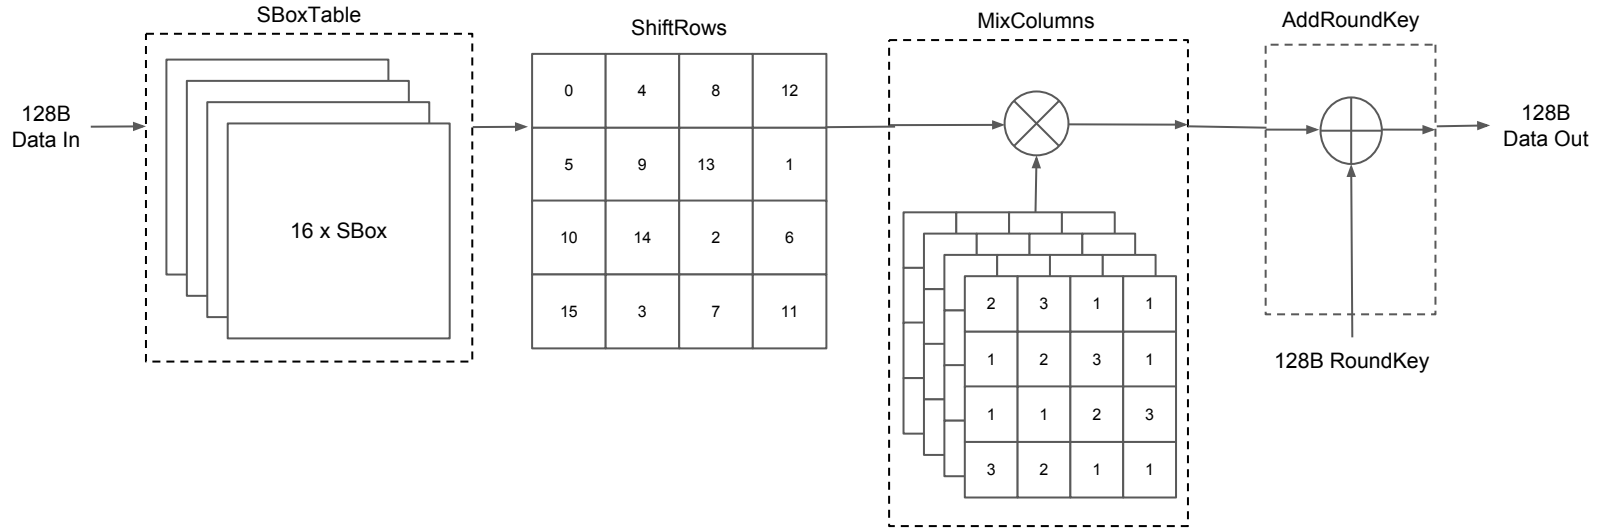


```
class PolyCompute(
    p: RSParams,
    numCells: Int,
    numInputs: Int,
    reverse: Bool
)
```



$(c_1 * x + c_2) x + \dots$

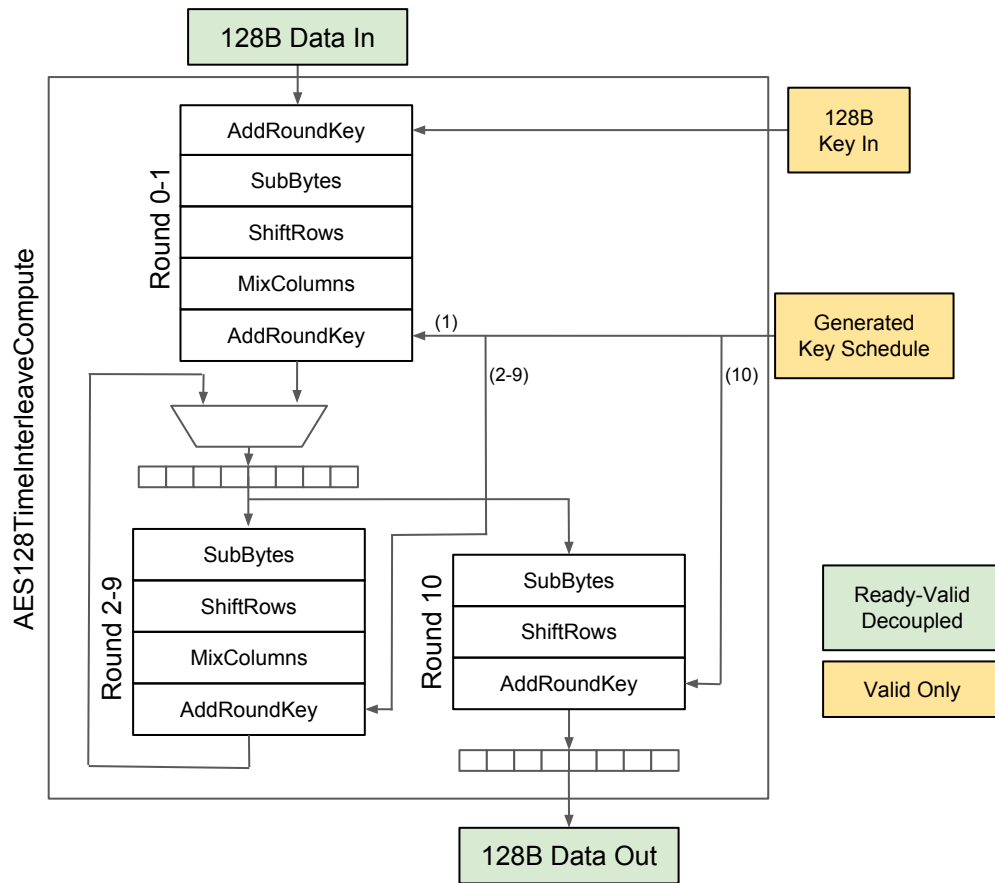
Encryption: AES Encryption



One AES Encrypt Stage
AES runs this 9 times, with a few extra substages before and after

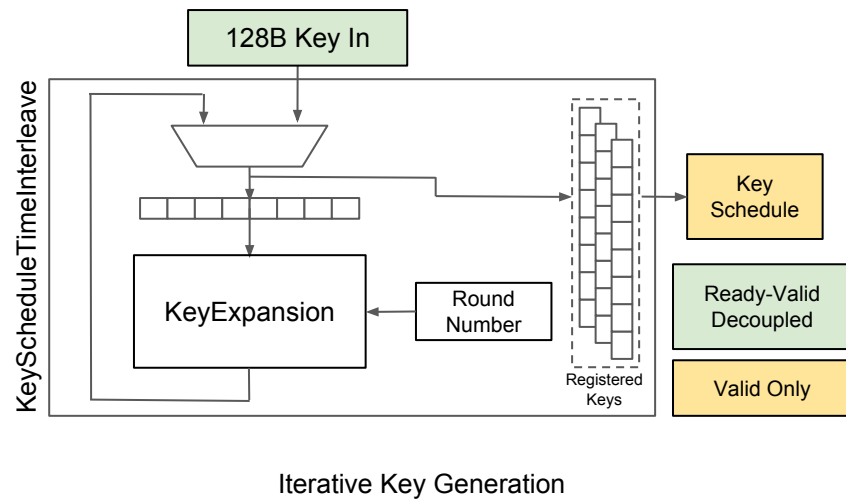
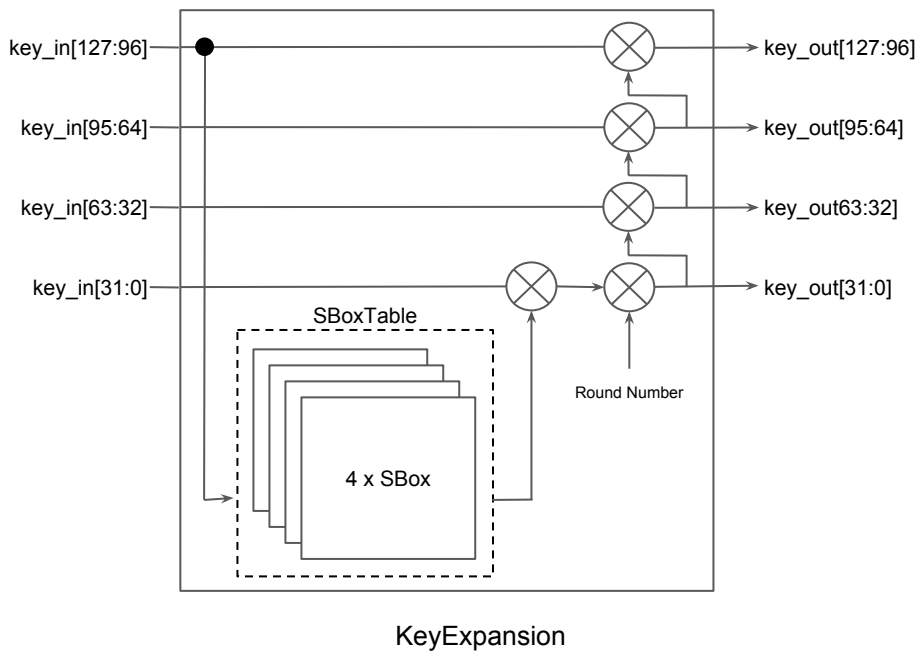
Encryption: AES 128 ECB

- Separate Compute blocks for encrypt and decrypt, with a shared key schedule generator
- Shared Key schedule generator reduces area and compute costs
- Iterative approach is easier to design and performance competitive with pipelined approached

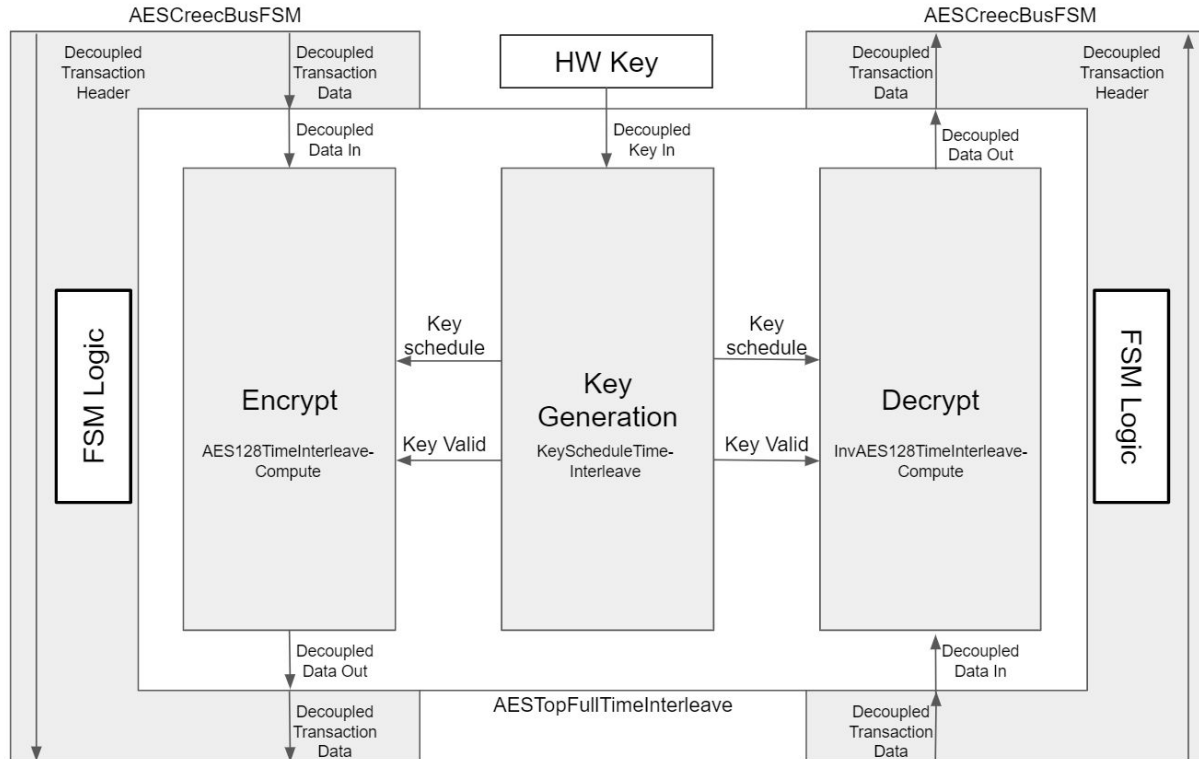


Encryption Block Diagram

Encryption: Key Generation



Encryption: Top Integration with CREECBus



Encryption Block Diagram

Encryption: Brief Results

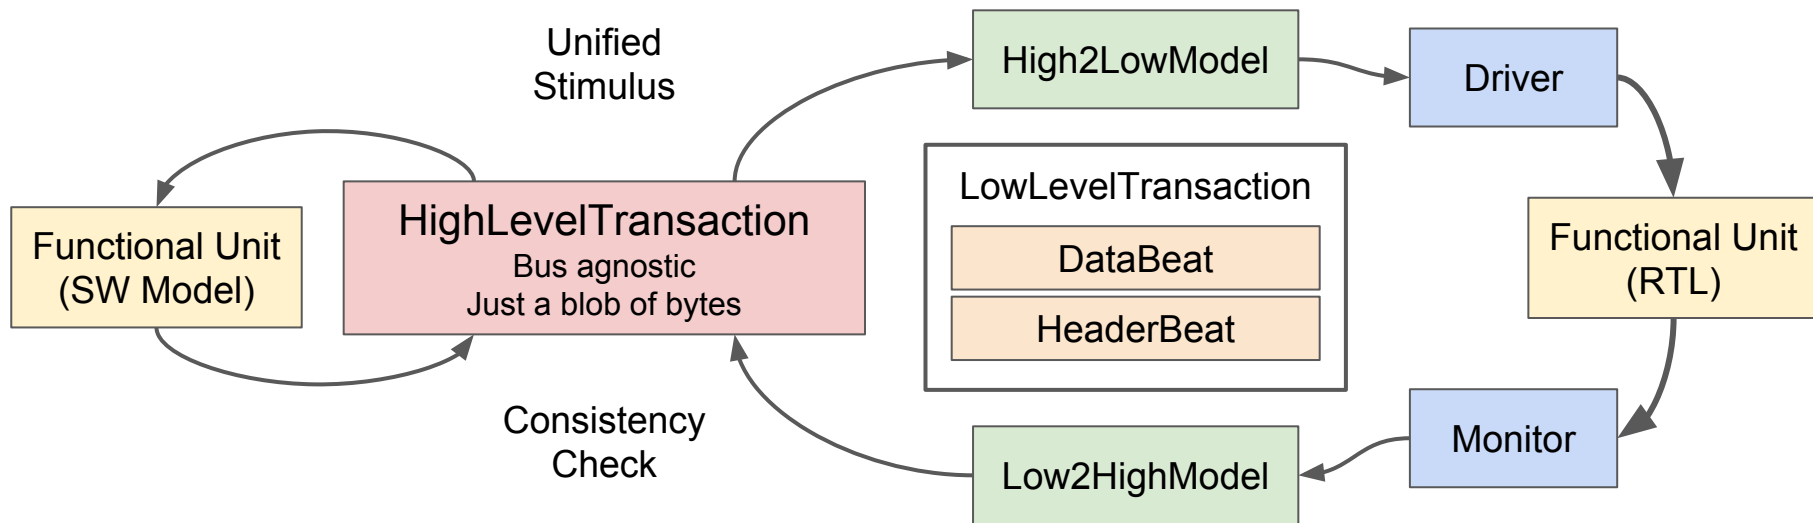
- Test latency and critical path are as expected
- Cipher area dominated by SBox Luts
- Key generation area dominated by key registers
- Future work: SBox reutilizations, exploring other operating modes, and time-interleaved performance

	Instance	Module	Cells
1	top		10781
2	AESTop	AESTopFullTimeInterleave	9360
3	keygen	KeyScheduleTimeInterleave	1858
4	rcon_gen	RCON	3
5	roundkey_gen	KeyExpansion	352
6	encrypt	AES128TimeInterleaveCompute	3545
7	stage0_addRoundKey	AddRoundKey_1	18
8	stage1_cipher	AESCipherStage_1	912
9	sub_byte	SubByte_2	640
10	mix_columns	MixColumns_1	144
11	MM0	MixColumnsMM_7	36
12	MM1	MixColumnsMM_6	36
13	MM2	MixColumnsMM_5	36
14	MM3	MixColumnsMM_4	36
15	add_round_key	AddRoundKey_7	128
16	AESStage	AESCipherStage	912
17	sub_byte	SubByte_1	640
18	mix_columns	MixColumns	144
19	MM0	MixColumnsMM_1	36
20	MM1	MixColumnsMM_2	36
21	MM2	MixColumnsMM_3	36
22	MM3	MixColumnsMM	36
23	add_round_key	AddRoundKey_2	128
24	stage10	AESCipherEndStage	768
25	sub_byte	SubByte	640
26	add_round_key	AddRoundKey_3	128
27	decrypt	InvAES128TimeInterleaveCompute	3957
28	stage0	InvAESCipherInitStage	768
29	add_round_key	AddRoundKey_4	128
30	inv_sub_byte	InvSubByte_1	640
31	InvAESStage	InvAESCipherStage_1	1064
32	add_round_key	AddRoundKey_6	128
33	inv_mix_columns	InvMixColumns_1	296
34	MM0	InvMixColumnsMM_7	74
35	MM1	InvMixColumnsMM_6	74
36	MM2	InvMixColumnsMM_5	74
37	MM3	InvMixColumnsMM_4	74
38	inv_sub_byte	InvSubByte_2	640
39	stage9	InvAESCipherStage	1064
40	add_round_key	AddRoundKey_5	128
41	inv_mix_columns	InvMixColumns	296
42	MM0	InvMixColumnsMM_1	74
43	MM1	InvMixColumnsMM_2	74
44	MM2	InvMixColumnsMM_3	74
45	MM3	InvMixColumnsMM	74
46	inv_sub_byte	InvSubByte	640
47	stage10	AddRoundKey	18
48	encrypt_FSM	AESCREECBusFSM_1	357
49	decrypt_FSM	AESCREECBusFSM	357

FPGA Synthesized usage (no flattening or optimization)

Software Modeling / RTL TLM

- Transaction-level modeling (TLM) enables unified stimulus to be used to drive and compare software and RTL algorithm implementations
- We designed a TLM framework in Scala and implemented SW models
- We integrated Chisel testers² to create CREECBus drivers and monitors



Transaction-Level Modeling

- **Transactions** are abstract; they represent a chunk of data (and/or control)
- **High-level** transactions are
 - Bus-**agnostic** (can work on AXI4-Stream or CREECBus)
 - Bus-parameterization **agnostic** (can be driven on a 64-bit/128-bit bus)

```
trait Transaction
```

```
abstract class CREECT extends Transaction
```

```
case class CREECHighT(data: Seq[Byte]) extends CREECT
```

- All our software models operate on high-level transactions
- RTL doesn't see high-level transactions directly

Transaction-Level Modeling

- **Low-level** transactions are
 - **Bus-specific** (represents bus-specific transactions)
 - Bus-parameterization **specific** (bus data width is fixed)

```
abstract class CREELowT extends CRECT
```

```
case class HeaderBeat(len: Int, id: Int)(p: BusParams) extends CREELowT {  
    require(len <= (p.maxBeats - 1))  
    require(id <= p.maxInFlight)  
}
```

```
case class DataBeat(data: Seq[Byte], id: Int)(p: BusParams) extends CREELowT {  
    require(data.length == p.bytesPerBeat)  
    require(id <= p.maxInFlight)  
}
```

Software Models

- SW models have 1 input and 1 output port and are untimed
 - Concrete models implement **process()**
- High2Low and Low2High are implemented as SW models

```
abstract class SoftwareModel[I <: Transaction, O <: Transaction] {  
  def process(in: I) : Seq[O]  
  def compose[O2 <: Transaction](s: SoftwareModel[O, O2]): SoftwareModel[I,  
O2]  
}
```

```
class EncryptModel extends SoftwareModel[CREECHighT, CREECHighT] {  
  override def process(in: CREECHighT) = {  
    val encryptedData = AESEncryption.encrypt(key, in.data)  
    Seq(in.copy(data = encryptedData))  
  }  
}
```

Software Model Composition

- Models can be composed to create data pipelines
- Models are tested by hand, then are used to produce golden reference output for RTL

```
val model =  
  CompressModel ->  
  PadderModel(16) ->  
  EncryptModel ->  
  ECCEncoderModel(RS16_8_8) ->  
  CommChannel(RS16_8_8, noiseByteLevel=4) ->  
  ECCDecoderModel(RS16_8_8) ->  
  DecryptModel ->  
  StripperModel ->  
  DeCompressModel  
val out = model.processTransactions(dataIn)  
assert(out == dataIn)
```

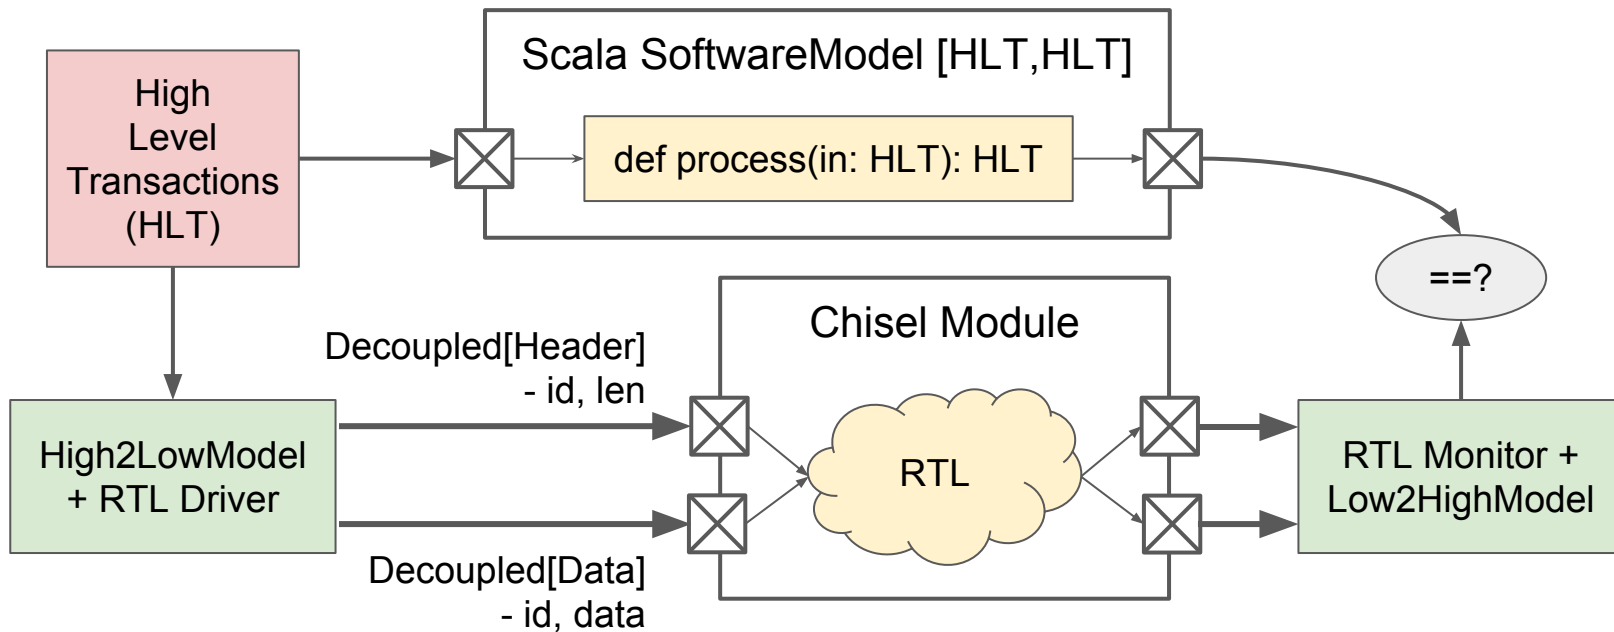
RTL Simulation and Verification

- **Driver/monitor** abstract away bus interaction from core test logic
- Test assertions operate on high-level transactions

```
val tx = CREECHighLevelTransaction(Seq(1, 2, 3, 4, 5, 6, 7, 8))
"the Compressor module" should "compress" in {
  test(CompressorRTL) { c =>
    val model = CompressorModel
    val outGold = model.processTransactions(tx)
    val driver = CREECDriver(c.io.in, c.clock)
    val monitor = CREECMonitor(c.io.out, c.clock)

    driver.pushTransactions(tx)
    c.clock.step(100)
    val out = monitor.receivedTransactions.getAll
    assert(out == outGold)
  }
}
```

Putting It All Together



- We used this scheme to test our entire CREECelerator pipeline

Results: Cycle Counts

- Cycle counts measured based on 512 bytes of ASCII text (one transaction)
- Full pipeline tests push data through each block in succession
 - The starred number is high because we corrupted the data before ECC decoding
- The standalone tests were run with each block in isolation

	Full Pipeline		Standalone Blocks	
	Encode	Decode	Encode	Decode
Compression	1764	1707	1764	1707
Encryption	474	474	460	460
ECC	1258	*7022	1220	2308
Total	3496	9203	3444	4475

- High numbers in compression are due to the bus semantics
 - Header length is not known until the end, but the header must be sent first

Results: Area (Vivado Synthesis with VC707)

- No Block RAMs or DSPs
- Compressor buffers entire transaction in memory

Module	LUTs	FFs
Compressor	29161	21403
Decompressor	29035	21431
ECCEncoderTop	160	362
ECCDecoderTop	3622	1374
AES128TimeInterleaveCompute	3238	274
InvAES128TimeInterleaveCompute	3453	263
KeyScheduleTimeInterleave	368	1421
AESCREECBusFSM	29	321
Total	69800	48757

Results: Critical Path (Vivado Synthesis with VC707)

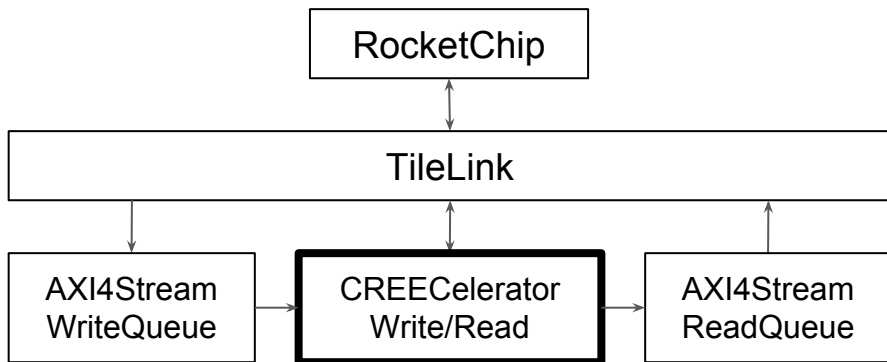
- Post-synthesis timing results
- All blocks were able to meet a target of 190 MHz
- The critical path in the decompressor can be cut at the cost of area
- Corresponds to a write/read throughput of 33.7 MB/s / 28.5 MB/s (very bad)

Top Module	Critical Path (ns)
Compressor	3.51
Decompressor	5.13
ECCEncoderTop	3.04
ECCDecoderTop	4.72
AESTopCREECBus	4.35

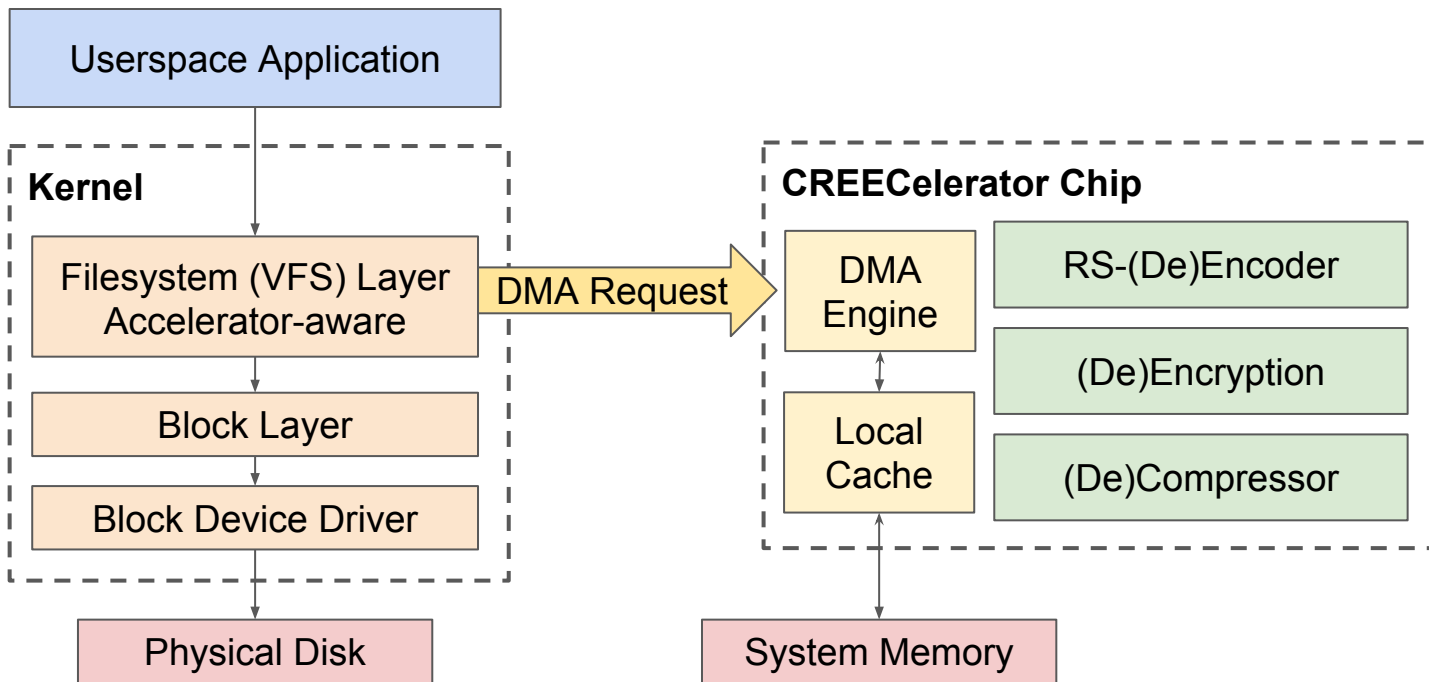
Rocket-chip Integration (it works!)

```
./simulator-freechips.rocketchip.system-CR  
EECWConfig ../tests/creec.riscv  
Sending data 578437695752307201  
...  
Received header: 8 1 1 1 4 8 0  
OUT: (i=0, j=0) creecW=31, gold=31  
OUT: (i=0, j=1) creecW=30, gold=30  
...  
OUT: (i=7, j=7) creecW=-99, gold=-99  
creecW PASSED!
```

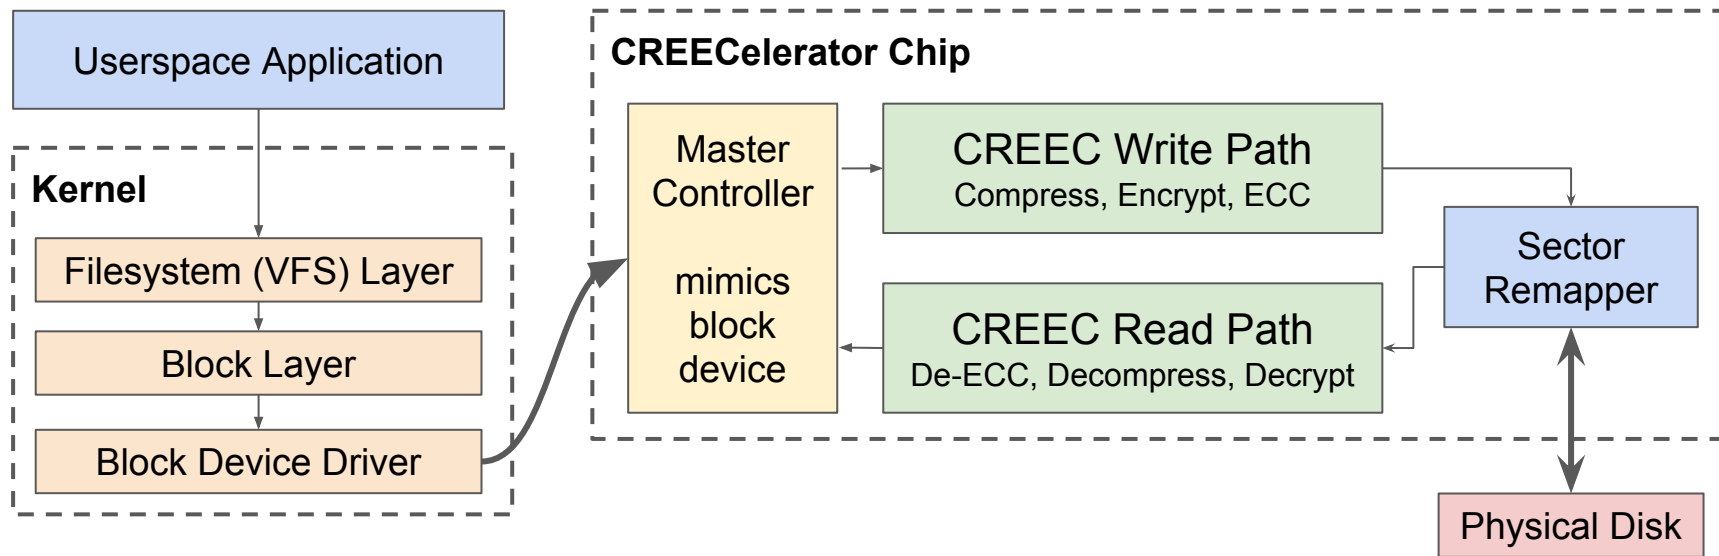
- The CREECelerator write and read paths were attached to Rocket just like CORDIC
- creec.c pushes data into the pipeline and reads the response



Application: Filesystem-Aware Accelerator

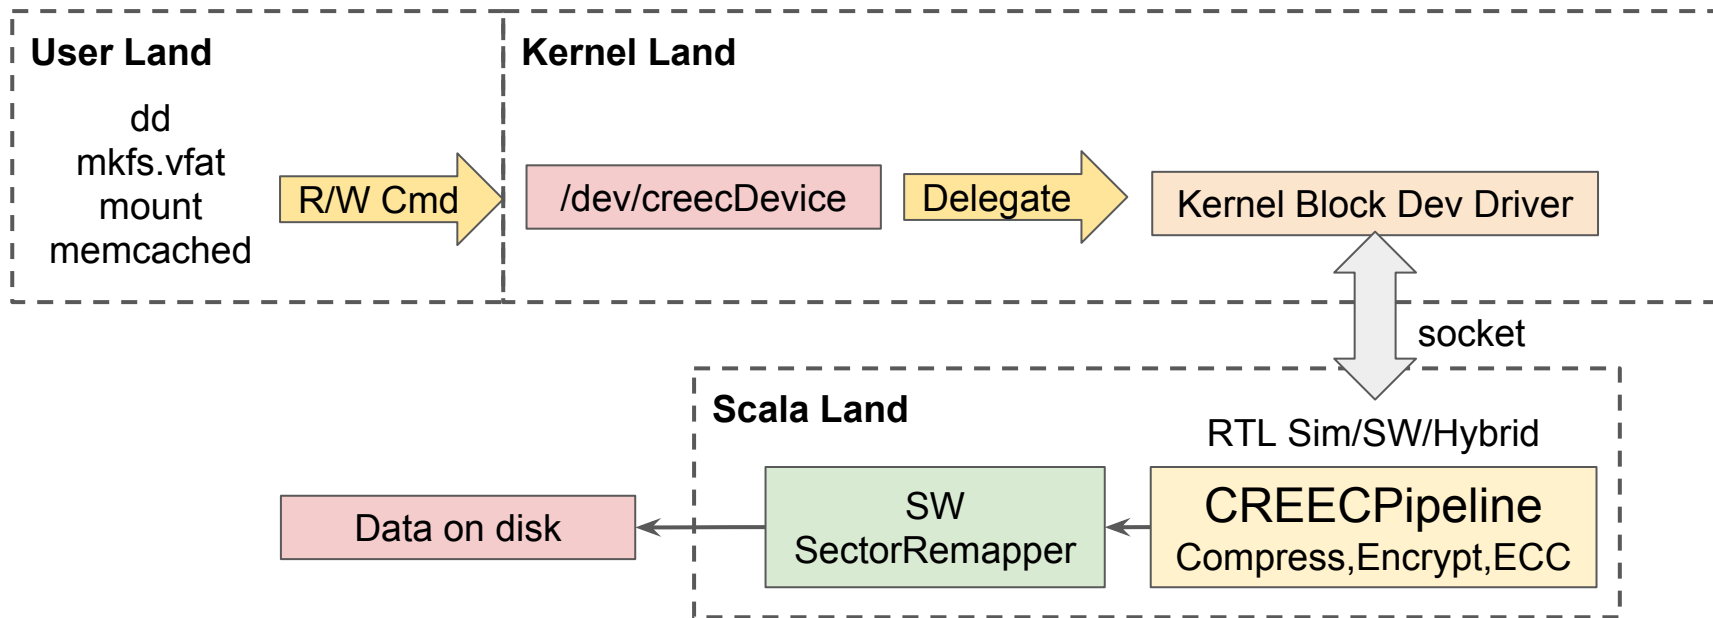


Application: Transparent Storage Utility Accelerator

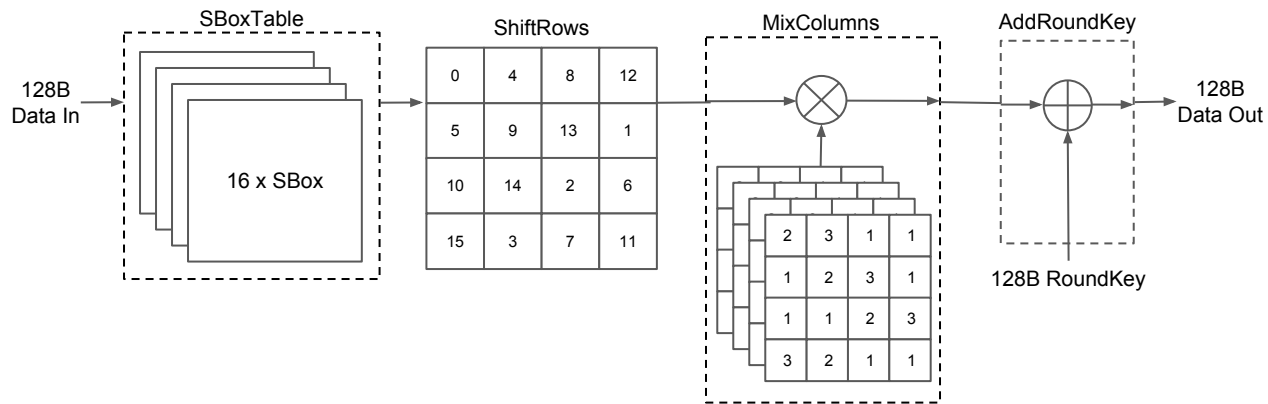
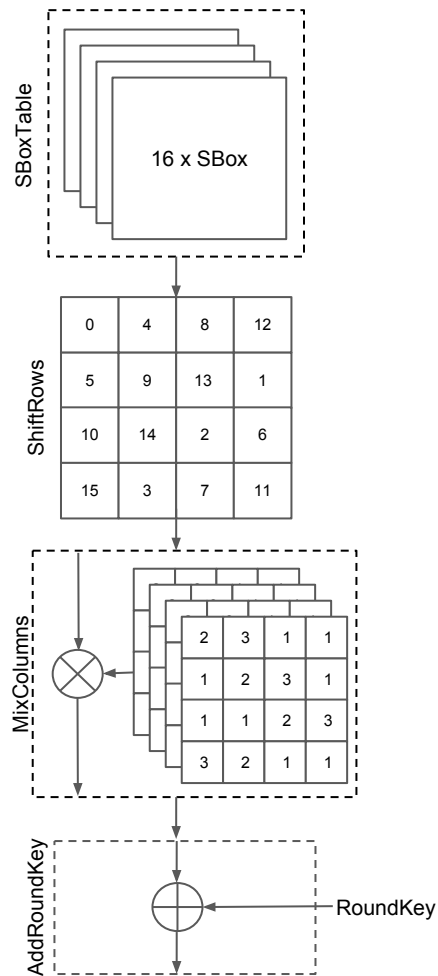


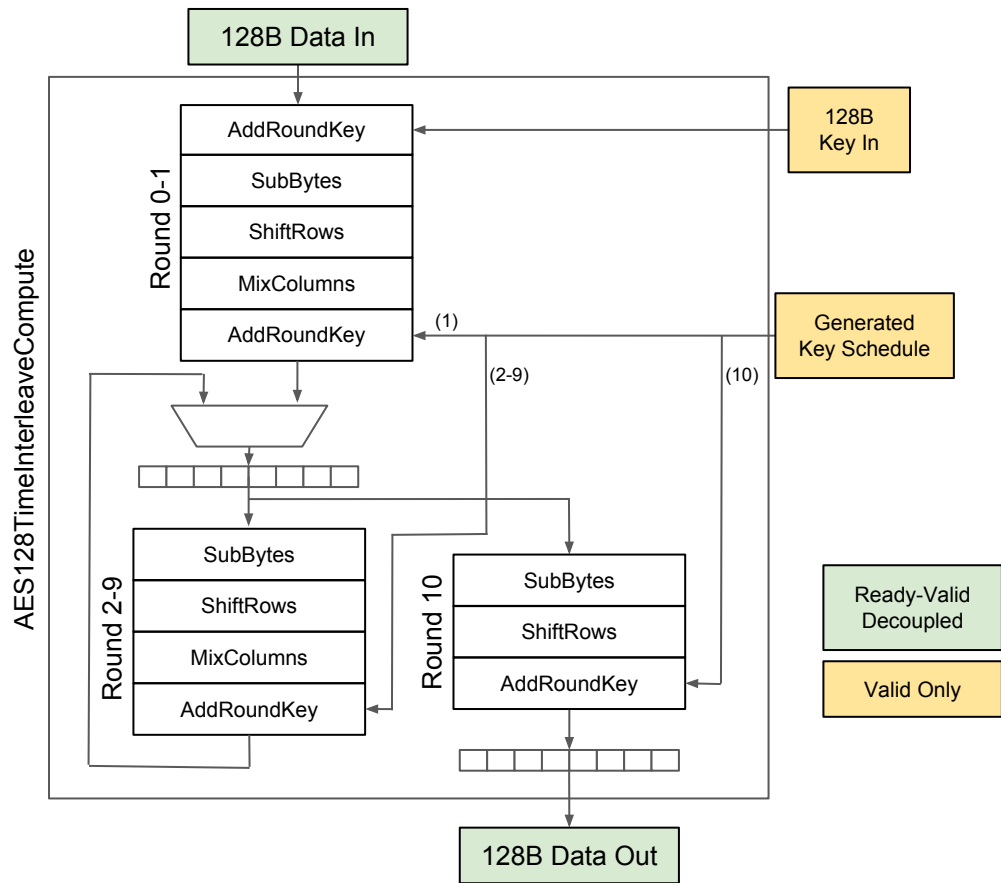
Future Work: System-Level Testing

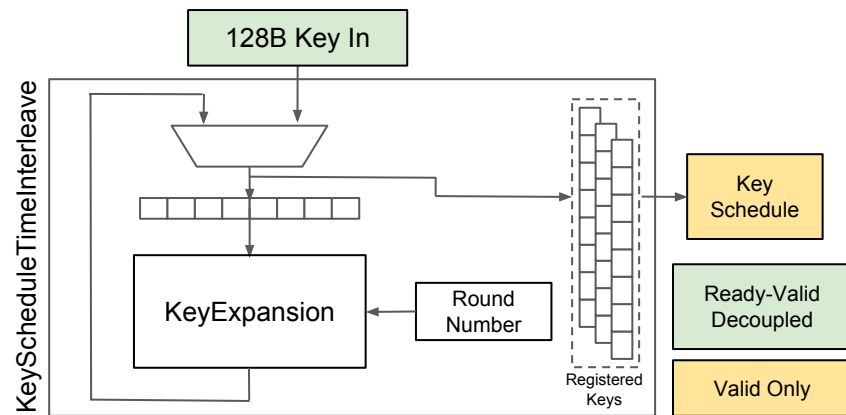
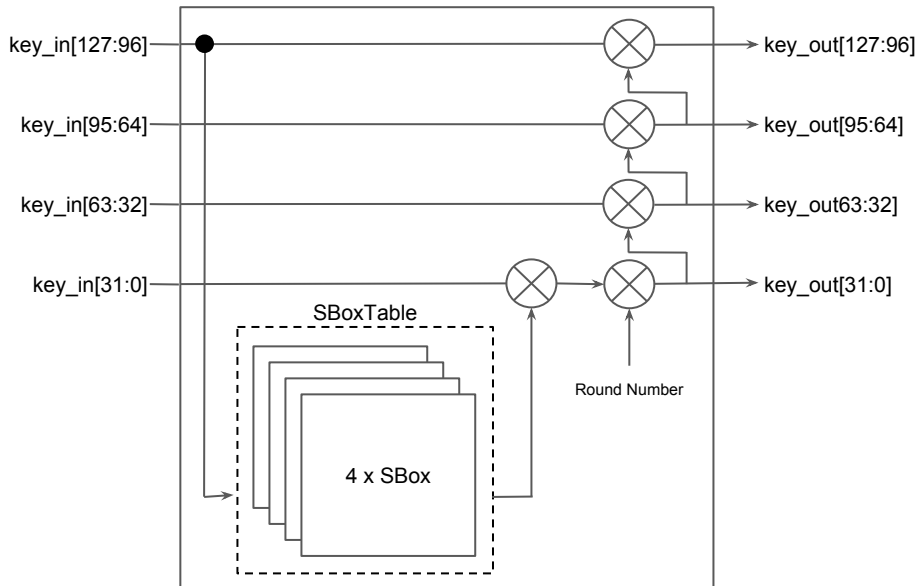
- Generate real block device traffic and funnel to software and RTL models of the CREEC pipeline; implement the SectorRemapper in software



Extra slides







AESTopFullTimeInterleave

